

# Machine Learning

MA842

# Course content

- Introduction to machine learning
- Decision Tree Learning
- Artificial Neural Networks
- Computational Learning Theory
- Instance-Based Learning
- Genetic Algorithms
- Analytical Learning
- Reinforcement Learning.

# Machine Learning : Definition

- "A computer program is said to **learn from experience E** (without being explicitly programmed) with respect to **some class of tasks T** and **performance measure P**, if its performance at tasks in T, as measured by **P**, **improves** with experience E."

$$E * T = P$$

- Tom Mitchell

- Example: playing checkers.
- E = the experience of playing many games of checkers
- T = the task of playing checkers.
- P = the probability that the program will win the next game.


# Application of ML

- **Transportation System:** Google maps provides the best route to reach the destination based on GPS location such that the congestion is minimum.  
Task – Build map of current traffic  
Experience – Learn the live scenario using GPS location, average velocity
- **Virtual Personal Assistants:** Collect and refine the information on the basis of our previous involvement with the personal assistant.
- **Product recommendations based on browsing history:** The shopping website or the app recommends some items that somehow matches with our taste.
- **3D printing of machine parts:** The sensor integrated system learns to minimize the production of defective machine parts by learning from previous flaws.

<https://www.youtube.com/watch?v=ahRcGObyEZo>


# Additional Examples

- Face detection in AI glasses for blind
- Handwriting recognition
- Autonomous vehicle
- Robot to clean the house
- Email filtering
- Weather prediction
- Medical diagnosis
- Stock market analysis



Source: <https://semiengineering.com/deep-learning-spreads/>


# Where does ML stand?



Source: <https://www.edureka.co/blog/ai-vs-machine-learning-vs-deep-learning/>

# Types of ML

- Supervised
  - Classification – A product is defective or not
  - Regression – Predict the price of a house
- Unsupervised
  - Association - People that buy X also tend to buy Y
  - Clustering - Grouping customers by purchasing behavior
- Reinforcement – autonomous driving car



Source: [https://en.wikipedia.org/wiki/Reinforcement\\_learning](https://en.wikipedia.org/wiki/Reinforcement_learning)

# Designing a learning system

- Choosing the training experience
  - Type of training experience : Direct / Indirect feedback (credit assignment problem)
  - Degree to which the learner controls the sequence of training examples.
  - Represent distribution of examples over which performance is measured.



- Choosing the target functions

- Determine exactly what type of knowledge will be learned and how this will be used by performance program.
- Target function  $V$ : Board  $\rightarrow$  Real number
- Define target value  $V(b)$  for an arbitrary board state  $b$ :

Final board state $b$	Value $V(b)$
Won	100
Lost	-100
Drawn	0

If  $b$  is not final state,  $V(b) = V(b')$ , where  $b'$  is best final board state

- Choosing the representation for target function

- Linear function:

$$\hat{V}(b) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_6$$

Where:

$w_0, w_1, w_2, w_3, w_4, w_5$ , and  $w_6$  - weights

$x_1$  : Number of black pieces

$x_2$  : Number of white pieces

$x_3$  : Number of black king pieces

$x_4$  : Number of white king pieces

$x_5$  : Number of black pieces threatened by white

$x_6$  : Number of white pieces threatened by black

In general,  $x_1$  to  $x_n$  represents the features projected on n dimensional feature space

- Choosing the function approximation algorithm

To learn  $\hat{V}$  we need set of training examples,  $\langle b, V_{train}(b) \rangle$

Example:  $\langle \langle x_1 = 3, x_2 = 0, x_3 = 1, x_4 = 0, x_5 = 0, x_6 = 0 \rangle, +100 \rangle$

**Aim:**

**To derive such training examples and adjust the weight values to best fit these training examples  $\{\langle b, V_{train}(b) \rangle\}$**

Rule for **estimating the training values** for any intermediate board state  $b$   
 $V_{train}(b) \leftarrow \hat{V}(\text{successor}(b))$

where  $\hat{V}$  is the learner's current approximation to  $V$  and  $\hat{V}(\text{successor}(b))$  denotes the next board state following  $b$

- Adjusting the weights

- Define the best hypothesis or set of weights such that squared error  $E$  between training values and the values predicted by the hypothesis  $\hat{V}$  is minimum.

$$E \equiv \sum_{\langle b, V_{train}(b) \rangle \in \text{training examples}} (V_{train}(b) - \hat{V}(b))^2$$

One algorithm to incrementally refine the weights as new training examples become available is least mean square (LMS) training rule.

# LMS Algorithm

## Weight update rule

- For each training example  $\langle b, V_{train}(b) \rangle$ 
  - Use the current weights to calculate  $\hat{V}(b)$
  - For each weight  $w_i$ , update it as
  - $w_i \leftarrow w_i + \eta (V_{train}(b) - \hat{V}(b)) x_i$

Where  $\eta$  is a constant that moderates the size of the weight update.

If  $(V_{train}(b) - \hat{V}(b)) = 0$ , no weights are changed

If  $(V_{train}(b) - \hat{V}(b)) > 0$ , each weight is increased in proportion to the value of its corresponding feature  $x_i$

# An example of LMS

Suppose we have a data set of 6 points as shown:

$i$	$x_i$	$y_i$
1	1.2	1.1
2	2.3	2.1
3	3.0	3.1
4	3.8	4.0
5	4.7	4.9
6	5.9	5.9

We find the best fitting line as follows.

We define the Mean Squared Error function. Let  $g(x) = mx + b$  represent a generic line. Its mean squared error in approximating the data is then

$$MSE = \frac{1}{6} \sum_{i=1}^6 (y_i - g(x_i))^2 = \frac{1}{6} \sum_{i=1}^6 (y_i - (mx_i + b))^2$$

Expanding this we have

$$\frac{1}{6} \sum_{i=1}^6 (y_i^2 - 2x_i y_i m - 2y_i b + x_i^2 m^2 + 2x_i m b + b^2)$$

This sum can be rewritten as

$$\begin{aligned}MSE &= \frac{1}{6} \sum_{i=1}^6 y_i^2 - \frac{1}{6} \sum_{i=1}^6 2x_i y_i m - \frac{1}{6} \sum_{i=1}^6 2y_i b + \frac{1}{6} \sum_{i=1}^6 x_i^2 m^2 + \frac{1}{6} \sum_{i=1}^6 2x_i m b + \frac{1}{6} \sum_{i=1}^6 b^2 \\&= \frac{1}{6} \sum_{i=1}^6 y_i^2 - \left( \frac{1}{6} 2 \sum_{i=1}^6 x_i y_i \right) m - \left( \frac{1}{6} 2 \sum_{i=1}^6 y_i \right) b + \left( \frac{1}{6} \sum_{i=1}^6 x_i^2 \right) m^2 + \left( \frac{1}{6} 2 \sum_{i=1}^6 x_i \right) m b + b^2\end{aligned}$$

We see that this is a function of our variables  $m$  and  $b$ , since  $x_i$  and  $y_i$  are simply numbers from our data set. Evaluating the sums in this last expression, we have

$$\sum_{i=1}^6 x_i = 20.9, \sum_{i=1}^6 y_i = 21.1, \sum_{i=1}^6 x_i y_i = 88.49, \sum_{i=1}^6 x_i^2 = 87.07, \sum_{i=1}^6 y_i^2 = 90.05.$$

so the the mean squared function can be written

$$MSE(m, b) = 15.00833 - 29.49667m - 7.03333b + 14.51167m^2 + 6.96667mb + b^2$$

To find the minimal possible value of  $MSE$  we find the partial derivatives of  $MSE$  and set them equal to zero:

$$\frac{\partial MSE}{\partial m} = -29.49667 + 29.02333m + 6.96667b = 0$$

$$\frac{\partial MSE}{\partial b} = -7.03333 + 6.96667m + 2b = 0$$

This second equation can be rearranged to give

$$b = \frac{7.03333 - 6.96667m}{2} = 3.51667 - 3.48333m$$

Plugging this into the other equation we have

$$0 = -29.49667 + 29.02333m + 9.63333(3.51667 - 3.48333m) = 4.38057 - 4.53273m$$

so that


$$m = 0.96643$$

and

$$b = 3.51667 - 3.48333(0.96643) = 0.15028.$$

Thus, the best fitting line is

$$y = 0.96643x + 0.15028.$$





# Another example for LMS

Suppose we want to predict the mileage of a car from its weight and age

Weight (x 100 lb) $x_1$	Age (years) $x_2$	Mileage
31.5	6	21
36.2	2	25
43.1	0	18
27.6	2	30

What we want: A function that can predict mileage using  $x_1$  and  $x_2$

**Assumption:** The output is a linear function of the inputs

$$\text{Mileage} = w_0 + w_1 x_1 + w_2 x_2$$

**Learning:** Using the training data to find the *best* possible value of **w**

**Prediction:** Given the values for  $x_1$ ,  $x_2$  for a new car, use the learned **w** to predict the **Mileage** for the new car

- Inputs are vectors:  $\mathbf{x} \in \mathbb{R}^d$
- Outputs are real numbers:  $y \in \mathbb{R}$

For simplicity, we will assume that  $x_1$  is always 1.

- We have a training set

$$D = \{ (\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m) \}$$

That is  $\mathbf{x} = [1 \ x_2 \ x_3 \ \dots \ x_d]^T$

This lets makes notation easier

- We want to approximate  $y$  as

$$\begin{aligned} y = f_{\mathbf{w}}(\mathbf{x}) &= w_1 + w_2 x_2 + \dots + w_n x_n \\ &= \mathbf{w}^T \mathbf{x} \end{aligned}$$

$\mathbf{w}$  is the learned weight vector in  $\mathbb{R}^d$

*Question:* How do we know which weight vector is the *best* one for a training set?

For an input  $(\mathbf{x}_i, y_i)$  in the training set, the *cost* of a mistake is

$$|y_i - \mathbf{w}^T \mathbf{x}_i|$$

Define the cost (or *loss*) for a particular weight vector  $\mathbf{w}$  to be


$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

Sum of squared costs over the training set


One strategy for learning: *Find the  $\mathbf{w}$  with least cost on this data*

$$\min_{\mathbf{w}} \frac{1}{2} \sum_{i=1}^m (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

Learning: minimizing mean squared error





Two dimensional input  
Predict using  $y = w_1 + w_2 x_2 + w_3 x_3$



# Gradient Descent Algorithm ~ Example


**Q: Minimize  $y = f(x) = x^2 - 4x + 2$  subject to  $-1 \leq x \leq 6$**

- In this optimization problem, our objective is a convex function with  $x \in \mathbb{R}$ .
- Define  $x_i := x_i - \Delta t \frac{\partial y}{\partial x}$  where  $\Delta t$  is timestamp (learning rate)




# Impact of learning rate

Too low




A small learning rate requires many updates before reaching the minimum point

Just right



The optimal learning rate swiftly reaches the minimum point

Too high




Too large of a learning rate causes drastic updates which lead to divergent behaviors

# Solve

Consider the data given in table below. Perform linear regression using gradient descent assuming learning rate is 0.01

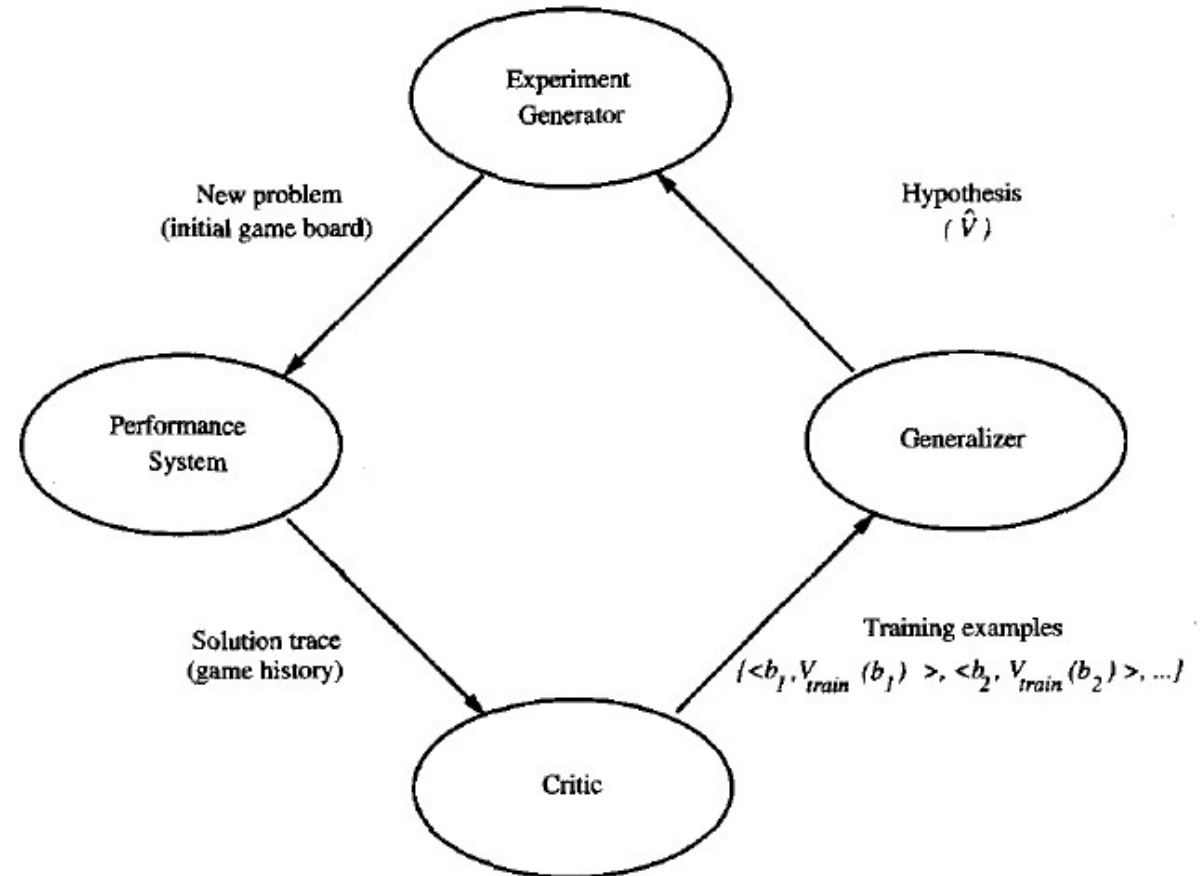
X No. of 'A' grades in I year	Y No. of 'A' grades in II year
1	1
2	3
4	3
3	2
5	5



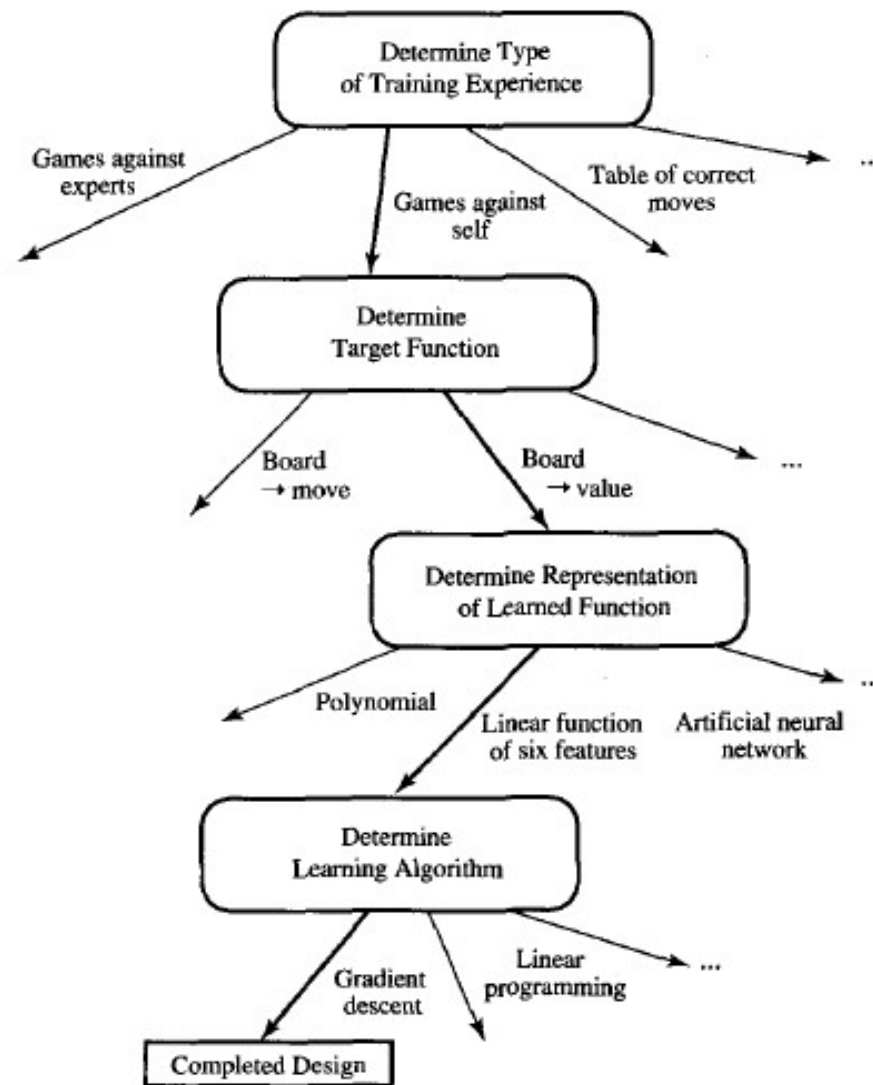


- Final Design

- Performance System
- Critic
- Generalizer
- Experiment Generator



# Design



# Perspectives in Machine learning

Machine learning involves searching a very large space of possible hypothesis to determine the one that best fits the observed data and any prior knowledge held by the learner.

Various representations are:

- Linear functions
- Decision trees
- Artificial Neural networks


# Issues in ML

- Amount of training data
- Missing and noisy feature values
- Appropriate feature selection
- Prior knowledge
- Strategy to choose next experience
- Learning rate
- Existence of algorithms to learn general target functions
- Overfitting and underfitting
- Bias variance effects

- **Overfitting:** A statistical model is said to be overfitted, when we train it with a lot of data. *That is, too much reliance on the training data.*
- **Underfitting:** A statistical model or a machine learning algorithm is said to have underfitting when it cannot capture the underlying trend of the data. *A failure to learn the relationships in the training data.*
- **High Variance:** model changes significantly based on training data.
- **High Bias:** assumptions about model lead to ignoring training data.
- Overfitting and underfitting cause poor **generalization** on the test set
- A **validation set** for model tuning can prevent under and overfitting

# Effects of overfitting and underfitting

- Consider a two class category problem to visualize the effects of underfitting and overfitting



# Avoiding overfitting issues

The commonly used methodologies are:

- **Cross- Validation:** A standard way to find out-of-sample prediction error is to use 5-fold cross validation.
- **Pruning:** Pruning is extensively used while building related models. It simply removes the nodes which add little predictive power for the problem in hand.
- **Regularization:** It introduces a cost term for bringing in more features with the objective function. Hence it tries to push the coefficients for many variables to zero and hence reduce cost term.

# Exercise

Pick a machine learning task. Identify T,E and P. Propose a target function to be learned and a target representation. Discuss the tradeoffs considered in formulating this task.